

Hiring Manager Interview Prep

Rahul Jain · Sr. Director IT Business Applications · Gusto
Enterprise Application AI Architect

1. KNOW YOUR INTERVIEWER

Rahul has a specific profile that shapes how he'll run this interview:

- Sr. Director IT Business Applications — he owns the enterprise systems stack (NetSuite, Zuora, Jira etc.) and is hiring someone to layer AI on top of it. His primary concern is: can you make my systems smarter without breaking them?
- Background: Infosys → Deloitte → Republic National Distributing → Freshworks → Gusto. Consulting + enterprise ops + SaaS. He evaluates candidates like a consultant: structured thinking, business impact, stakeholder alignment.
- Executive MBA in Finance, Data Science, ML & AI (Michigan Ross). He speaks both languages — business ROI and technical architecture. Don't over-index on either alone.
- He designed this role as a new headcount. He has a vision for it. He's looking for someone who can extend his thinking, not just execute tasks.

Tip: Rahul will likely drive toward business outcomes early. Lead every answer with the problem being solved and the impact delivered, then explain how you solved it — not the other way around.

2. OPENING SCRIPT — "TELL ME ABOUT YOURSELF"

Aim for 90 seconds. Practice this until it's natural.

"I'm a solutions architect with just over ten years of experience building enterprise-scale data and AI platforms. Most recently at Verizon, I've been the primary technical owner for a set of platforms that are directly relevant to what you're building at Gusto. The most relevant is probably the Agentic AI initiative I led — where I integrated LLMs into operational workflows using LangChain and LangGraph, connecting them to enterprise systems via production API automation. That work compressed process cycles from weeks to near-real-time and reclaimed thousands of analyst hours annually. It's real AI in production, not a prototype. I also architected a B2B Billing platform serving over two thousand enterprise clients, with full end-to-end data traceability from ingestion through auditing — which is directly analogous to what you'd need around something like Zuora. And a Supply Chain data platform on AWS that generated over twenty million dollars in working capital savings. What draws me to this role specifically is that I want to go deeper into the enterprise applications layer — not just build analytics on top of systems like NetSuite and Zuora, but actually wire AI into the workflows those systems run. That's the problem this role is designed to solve, and it's where I want to focus next."

3. THE ENTERPRISE SYSTEMS GAP — EXACT LANGUAGE

The recruiter flagged this twice. Rahul will probe it. Here is the full script for when he does.

Do NOT say: "I'm a fast learner" or "I can pick it up quickly." That sounds defensive. Do NOT over-explain or apologize. State it once, bridge immediately, move on.

WHEN HE ASKS: "HAVE YOU WORKED DIRECTLY WITH NETSUITE, ZUORA, OR JIRA?"

"Not as a developer or admin inside those specific systems — and I want to be straightforward about that. But I've built the equivalent of what they do at a similar scale. On the billing side: I architected a B2B Billing platform at Verizon serving over two thousand enterprise clients, with the full pipeline — ingestion, transformation, validation, and auditing. The core challenge of billing accuracy, dispute resolution, and financial traceability is the same problem Zuora is built to solve. I understand that domain deeply. On the ERP side: the Supply Chain platform I built integrates inventory, operational, and financial data flows across major OEMs and 3PL partners — that cross-functional data visibility is what NetSuite delivers for ERP. Different system, same architectural problem. For integration work like Jira: connecting operational and engineering systems via RESTful APIs is exactly what I've been doing. Wiring an AI agent to Jira's API is an integration problem, and integration is where I spend most of

my time. The part of this role I'm not doing yet is configuring those systems from the inside — but that's not what this role needs. This role needs someone to design the AI and integration layer on top of those systems. That's what I do."

If he pushes further: "What would your first 30 days look like in terms of getting up to speed on our systems?" Answer: shadow the NetSuite and Zuora admins for two weeks, map the data schemas and API surfaces, then identify the three highest-value integration points for AI agents. Don't start building until you understand what's running.

4. LIKELY QUESTIONS — FULL ANSWER SCRIPTS

Q1: Walk me through how you've turned a vague business request into an AI workflow.

Why he asks: He has Finance and Ops stakeholders who will come to you with fuzzy problems. He needs to know you can do discovery, not just build.

"The clearest example is the Agentic AI initiative at Verizon. The business ask came in as: 'analysts are spending weeks on these operational cycle reviews — can we do something about it?' That's not a technical spec, that's a pain point. My first step was to sit with the analysts and map the actual process — what data they were pulling, from which systems, what decisions they were making at each step, and where the bottlenecks were. What I found was that about seventy percent of the work was mechanical: pulling data from APIs, formatting it, applying a set of known rules, and escalating exceptions. That's automatable. I then designed an agentic workflow using LangChain and LangGraph — where each agent handled a discrete step: data retrieval, normalization, rule application, and exception flagging. The agents communicated through a structured context protocol so the outputs of one fed the inputs of the next reliably. We went from a multi-week cycle to near-real-time processing. The analysts shifted from doing the work to reviewing exceptions flagged by the agents. That's the outcome: not just automation, but a different operating model for that team."

Q2: How do you handle governance and auditability in AI systems?

Why he asks: He's responsible for systems that touch Finance and HR. One compliance failure is a serious incident. He needs to know you take this as seriously as he does.

"I treat governance as a design constraint, not a phase. It has to be baked in from the start because retrofitting audit controls into a live system is painful and error-prone. The clearest example is the Billing platform at Verizon. Every data state was tracked — ingestion timestamp, transformation applied, validation result, who or what triggered each step. If a dispute came in, we could reconstruct the exact data state that produced any given output. That's not just good engineering — in a financial context, that's a compliance requirement. For AI agents specifically, I add a few additional layers: every agent action is logged with its inputs, the model's output, and the decision made. Agents that touch financial or sensitive data require a human-in-the-loop approval step for anything above a defined confidence threshold. And I design for explainability — if someone asks why the agent made a decision, the audit log has the answer. I haven't worked in a formally SOX-audited environment, but the control patterns I build — immutable audit logs, access control, data lineage, exception escalation — are the same controls SOX requires. I'd map those patterns to Gusto's specific compliance framework in the first few weeks."

Q3: How do you measure whether an AI workflow is actually working?

Why he asks: His MBA is in Finance and Data Science. He cares about ROI and will want to know if you can quantify your own impact.

"I start by defining the baseline before anything is built — what is the current process, how long does it take, what's the error rate, what does it cost in people time. Without that baseline, you can't measure improvement. At Verizon I tracked four things for the Agentic AI initiative: process cycle time before and after, analyst hours per task, exception escalation rate — which tells you how often the agent is uncertain and needs human review — and agent accuracy on a held-out sample we reviewed manually each month. The Supply Chain platform had a cleaner financial metric: working capital savings from improved inventory data visibility. That came out to over twenty million dollars annually, which made the business case for continued investment obvious. What I've learned is that efficiency metrics alone aren't enough. You also need to track agent reliability over time — model drift is real, and an agent that was ninety percent accurate at launch can degrade without anyone noticing if you're not watching it. So I build monitoring in from day one, not as an afterthought."

Q4: This is a new role with no playbook. How do you approach building something from scratch?

Why he asks: He designed this role. He wants someone who can operate in ambiguity and build the function, not someone who needs clear requirements handed to them.

"I've built from scratch before. At Verizon, the Agentic AI initiative didn't exist when I started it — I identified the opportunity, made the case to leadership, designed the approach, and drove adoption across teams that had never worked with AI agents before. My approach is: spend the first few weeks in discovery mode before writing a single line of code. Talk to Finance, talk to Operations, talk to the system admins who know where the pain is. Map the highest-friction workflows. Then prioritize by the intersection of business impact and technical feasibility — you want quick wins that build credibility alongside the more complex work. The second thing I do early is establish the governance and documentation framework — reusable patterns, integration design standards, decision logs. That might sound like overhead, but it's what lets you move fast without creating technical debt that bites you six months later. I'm comfortable in ambiguity. What I need is access to stakeholders and clear alignment on what 'success in year one' looks like from your perspective — which is actually a question I'd like to ask you."

Q5: How do you work with Finance and Operations stakeholders who aren't technical?

Why he asks: He came from consulting (Deloitte) where stakeholder management is everything. He will not hire someone who can only work with engineers.

"The key for me is translating into the language the stakeholder already speaks — which is usually time and money. When I was working with the operations team at Verizon on the Agentic AI initiative, I didn't lead with 'here's how LangGraph works.' I led with: 'your analysts are spending three weeks per cycle on work that takes the system thirty seconds. Here's what that costs in headcount and here's what we can do about it.' That's a conversation a Finance leader will engage with immediately. Where I've found non-technical stakeholders are most valuable is in defining what 'right' looks like. They know the edge cases, the exceptions, the situations where automated decisions would be wrong and a human has to stay in the loop. I involve them in that design early — not just in requirements gathering, but in reviewing the agent behavior during testing. They catch things that no engineer would catch. The trap I try to avoid is over-promising. AI systems have confidence thresholds and failure modes, and if the stakeholder expects a hundred percent accuracy and gets ninety-five, that's a trust problem. I set expectations around what the system will do well and what it will escalate."

Q6: What's your experience with MCP — Model Context Protocol?

Why he asks: MCP is central to the job description. He may probe whether you've used it directly.

"I haven't implemented MCP using that specific protocol standard by name, but the architectural pattern it implements is what I've been building. The core idea of MCP is a standardized way to expose enterprise system context — data, tools, and state — to AI agents securely and reliably. That's exactly what I built at Verizon. My agentic workflows used a structured context-passing layer between agents and the enterprise systems they connected to — defining what data each agent could access, in what format, with what access controls, and with full logging of what was retrieved and used. That's the same problem MCP solves, even if I implemented it through LangGraph's state management and RESTful API integration patterns rather than the formal MCP spec. I've been following MCP's development and I understand the protocol — it's essentially a standardized tool-calling and context injection layer. Adopting the formal spec on top of my existing integration architecture would be a relatively small step."

5. QUESTIONS TO ASK RAHUL

Ask 3–4 of these. The ones marked with ★ are highest priority.

★"What are the one or two enterprise workflows where you feel the pain most acutely right now — where if AI could help, it would have an immediate impact?"

Why it works: Shows you're already thinking about prioritization and business value, not just technology. His answer also tells you exactly where to focus.

★"How much autonomy will this role have to go directly to Finance and Operations stakeholders — or does everything flow through the IT team?"

Why it works: This is a practical question about how you'll actually do the job. It also signals that you know cross-functional access matters for this kind of work. His answer will tell you a lot about how the function operates.

"When you imagine this role at 12 months — what does it look like if it's gone really well?"

Why it works: Gets him to articulate his vision, which is what you're being hired to execute. It also reframes the conversation toward the future rather than your past experience.

"How does Gusto currently handle AI governance — is there a central framework, or is it team-by-team?"

Why it works: Shows you care about compliance and aren't going to go rogue. Also genuinely useful to know — the answer shapes how you'll work.

"What's your read on how Finance and Operations leadership feels about AI automation right now — enthusiastic, skeptical, somewhere in between?"

Why it works: Stakeholder buy-in is make-or-break for this role. This tells you what kind of change management work is ahead, and it's a question he'll appreciate because it shows you've thought beyond the technical build.

One thing to remember: Rahul built this role from scratch. He wants to hire someone who makes it bigger than he imagined, not someone who just fills the spec. Show him you have a point of view on where enterprise AI is going — and that you're already living it.